



Xi'an Jiaotong-Liverpool University

西交利物浦大學

SQL Select III

Jianjun Chen (Jianjun.Chen@xjtlu.edu.cn)

Content

- Index
- Set Operations
- Missing information:
 - Dealing with nulls.
 - Making use of default values.

Index

- What are indices and why are they needed:
- <https://www.youtube.com/watch?v=fsG1XaZEa78>

Table structure

Relation view

#	Name	Type	Collation	Attributes	Null	Default	Comments	Extra
<input type="checkbox"/>	1	game_id	int(11)		No	None		
<input type="checkbox"/>	2	game_title	varchar(200) utf8mb4_general_ci		No	None		
<input type="checkbox"/>	3	game_release_date	date		Yes	NULL		
<input type="checkbox"/>	4	age_rating	int(11)		No	None		
<input type="checkbox"/>	5	developer_id	int(11)		No	None		

☐ Check all

With selected:

☐ Browse

☐ Change

☐ Drop

☐ Primary

☐ Unique

Print

Propose table structure

Track table

Move columns

Normalize

Add

1

column(s)

after developer_id

Go

Indexes

Action	Keyname	Type	Unique	Packed	Column	Cardinality	Collation	Null	Comment
<input type="checkbox"/> Edit	<input type="checkbox"/> Drop	PRIMARY	BTREE	Yes	No	game_id	0	A	No

Create an index on

1

columns

Go

Index


- An index helps to speed up select queries and where clauses.
- but it slows down data input, with the update and the insert statements.
- Primary key and unique key affects the availability of certain functions of SQL.
- For example:
 - `_rowid` is a special column that is only accessible when you have created a primary key or a unique key.

Set Operations


UNION, INTERSECT and EXCEPT

SET operations

- UNION, INTERSECT and EXCEPT
 - Treat the tables as sets and are the usual set operators of union, intersection and difference
 - Only UNION is supported in MySQL. The other two can be simulated with subqueries.
- They all combine the results from **two select statements**
- The results of the two selects should have the same columns and corresponding data types
 - Union compatible.



A EXCEPT B



A INTERSECT B



A UNION B

UNION: Example

- Find, in a single query, the average mark for each student and the average mark overall.

Grades		
Name	Code	Mark
Jane	IAI	54
John	DBS	56
John	IAI	72
James	PR1	43
James	PR2	35
Mary	DBS	60

UNION

1. The average for each student:

```
SELECT Name, AVG(Mark) AS Average  
      FROM Grades  
      GROUP BY Name;
```

2. The average overall:

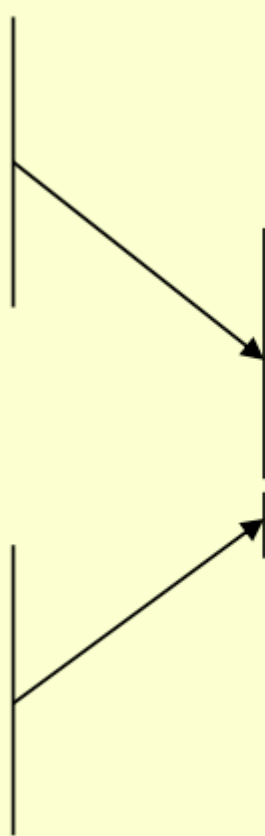
```
SELECT 'Total' AS Name,  
      AVG(Mark) AS Average  
      FROM Grades;
```

UNION

```
SELECT Name ,  
        AVG (Mark) AS Average  
FROM Grades  
GROUP BY Name
```

UNION

```
SELECT  
    'Total' AS Name ,  
    AVG (Mark) AS Average  
FROM Grades ;
```



Name	Average
James	39
Jane	54
John	64
Mary	60
Total	53.3333

Missing Information

Dealing with nulls in SQL

Making use of default values

Missing Information

- Sometimes we don't know what value an entry in a relation should have.
 - Case 1: We know that there is a value, but don't know what it is.
 - Case 2: There is no value at all that makes any sense.
- Two main methods have been proposed to deal with this
 - `NULLS` can be used as markers to show that information is missing .
 - A default value can be used to represent the missing value.

NULL

- **NULL** Represents a state for an attribute that is currently unknown or is not applicable for this tuple.
 - **NULLs** are a way to deal with incomplete or exceptional data.
 - **NULL** is a placeholder for missing or unknown value of an attribute.
It is not itself a value.
 - E.g. A new staff is just added, but hasn't been decided which branch he belongs to.
- Codd proposed to distinguish two types of **NULLs**:
 - **A-marks**: data Applicable but not known (for example, someone's age)
 - **I-marks**: data is Inapplicable (telephone number for someone who does not have a telephone, or spouse's name for someone who is not married)

Problems with NULLs

- Problems extending relational algebra operations to NULLs:
 - Selection operation: if we check tuples for “Mark > 40” and for some tuple Mark is NULL, do we include it?
 - Comparing tuples in two relations: are two tuples (with NULLs) and the same or not?
- Additional problems for SQL:
 - NULLs treated as duplicates?
 - Inclusion of NULLs in count, sum, average? If yes, how?
 - Arithmetic operations behaviours with argument NULL?

Theoretical Solutions

- Use three-valued logic instead of classical two-valued logic to evaluate conditions:
 - When there are no NULLs around, conditions evaluate to **true** or **false**, but if a NULL is involved, a condition might evaluate to the third value ('undefined', or '**unknown**')

a	b	a OR b	a AND b	a == b
True	True	True	True	True
True	False	True	False	False
True	Unknown	True	Unknown	Unknown
False	True	True	False	False
False	False	False	False	True
False	Unknown	Unknown	False	Unknown
Unknown	True	True	Unknown	Unknown
Unknown	False	Unknown	False	Unknown
Unknown	Unknown	Unknown	Unknown	Unknown

SQL NULLs in Conditions

Name	Salary
John	25,000
Mark	15,000
Anne	20,000
Chris	NULL

```
SELECT * FROM Employee  
Where Salary > 15,000;
```



Name	Salary
John	25,000
Anne	20,000

WHERE clause of SQL SELECT uses three-valued logic: only tuples where the condition evaluates to true are returned.

Salary > 15,000 evaluates to 'unknown' on the last tuple – not included

SQL NULLs in Conditions

Employee	
Name	Salary
John	25,000
Mark	15,000
Anne	20,000
Chris	NULL

```
SELECT * FROM Employee
      Where Salary > 15,000
      OR Name = 'Chris';
```



Name	Salary
John	25,000
Anne	20,000
Chris	NULL

Salary > 15,000 OR Name = 'Chris' is essentially
Unknown OR TRUE on the last tuple:

a	b	a OR b
Unknown	True	True

SQL NULLs in Arithmetic

Employee	
Name	Salary
John	25,000
Mark	15,000
Anne	20,000
Chris	NULL

```
SELECT Name,  
        Salary * 0.05 AS Bonus  
FROM Employee;
```



Name	Bonus
John	1,250
Mark	750
Anne	1,000
Chris	NULL

Arithmetic operations applied to NULLs result in NULLS

SQL NULLs in Aggregation


Employee	
Name	Salary
John	25,000
Mark	15,000
Anne	20,000
Chris	NULL

```
SELECT  
    AVG (Salary) AS Average,  
    COUNT (Salary) AS Count,  
    SUM (Salary) AS Sum  
FROM Employee;
```

Average = 20,000

Count = 3

Sum = 60,000



Using `COUNT (*)` would give 4, even if the name of Chris is changed to NULL.

SQL NULLs in GROUP BY

Employee	
Name	Salary
John	25,000
Mark	15,000
Anne	20,000
Jack	NULL
Sam	20,000
Chris	NULL

```
SELECT Salary,  
        COUNT(Name) AS Count  
FROM Employee  
GROUP BY Salary;
```



Salary	Count
NULL	2
15,000	1
20,000	2
25,000	1

NULLs are treated as equivalents in
GROUP BY clauses

SQL NULLs in ORDER BY

Employee	
Name	Salary
John	25,000
Mark	15,000
Anne	20,000
Jack	NULL
Sam	20,000
Chris	NULL

```
SELECT *  
FROM Employee  
ORDER BY Salary;
```



Employee	
Name	Salary
Chris	NULL
Jack	NULL
Mark	15,000
Anne	20,000
Sam	20,000
John	25,000

NULLs are considered and reported in ORDER BY clauses

Default Values

- Default values are an alternative to the use of `NULL`s
 - You can choose a value that makes no sense in normal circumstances.
 - These are actual values

```
age INT DEFAULT -1,
```

- Default values can have more meaning than `NULL`s. The followings are example default values of `VARCHAR`.
 - 'none'
 - 'unknown'
 - 'not supplied'
 - 'not applicable'
- Not all defaults represent missing information. It depends on the situation

Default Values: Example

- Default values are
 - “Unknown” for Name
 - -1 for Weight and Quantity
- -1 is used for Wgt and Qty as it is not sensible otherwise
- There are still problems:

```
UPDATE Parts SET  
Quantity = Quantity + 5
```

Parts			
ID	Name	Weight	Quantity
1	Nut	10	20
2	Bolt	15	-1
3	Nail	3	100
4	Pin	-1	30
5	Unknown	20	20
6	Screw	-1	-1
7	Brace	150	0

SQL Support

- SQL allows both `NULL`s and defaults:
 - A table to hold data on employees
 - All employees have a name
 - All employees have a salary (default 10000)
 - Some employees have phone numbers, if not we use `NULL`s

```
CREATE TABLE Employee  
(  
    Name VARCHAR(50) NOT NULL,  
    Salary INT DEFAULT 10000 NOT NULL,  
    Phone VARCHAR(15) NULL  
);
```


SQL Support

- SQL allows you to insert NULLs:

```
INSERT INTO Employee VALUES  
    ('John', 12000, NULL);  
UPDATE Employee SET Phone = NULL  
    WHERE Name = 'Mark';
```

- You can also check for NULLs:

```
SELECT Name FROM Employee  
    WHERE Phone IS NULL;  
SELECT Name FROM Employee  
    WHERE Phone IS NOT NULL;
```

SQL: The Final Example

Increasing the difficulty

The Final Example

- Examiners' reports
 - We want a list of students and their average mark
 - For first and second years the average is for that year
 - For finalists (third year) it is 40% of the second year plus 60% of the final year averages
- We want the results:
 - Sorted by year (desc), then by average mark (high to low) then by last name, first name and finally ID
 - To take into account of the number of credits each module is worth
 - Produced by a single query

Student			
ID	First	Last	Year

Grade			
ID	Code	Mark	YearTaken

Module		
Code	Title	Credits

Example Output

Student			
ID	First	Last	Year

Grade			
ID	Code	Mark	YearTaken

Module		
Code	Title	Credits



Year	Student.ID	Last	First	AverageMark
3	11014456	Andrews	John	81
3	11013891	Smith	Mary	78
3	11014012	Brown	Amy	76
3	11013204	Jones	Steven	76
3	11014919	Robinson	Paul	74
3	11013704	Edwards	Robert	73
⋮				
1	11027871	Green	Michael	45
1	11024298	Hall	David	43
1	11024826	Wood	James	40
1	11027621	Clarke	Stewart	39
1	11024978	Wilson	Sarah	36
1	11026563	Taylor	Matthew	34
1	11027625	Williams	Paul	31

Getting Started

- Finalists should be treated differently to other years
 - Write one SELECT for the finalists
 - Write a second SELECT for the first and second years
 - Merge the results using a UNION

QUERY FOR FINALISTS

UNION

QUERY FOR OTHERS

Table Joins

- Both subqueries need information from all the tables
 - The student ID, name and year
 - The marks for each module and the year taken
 - The number of credits for each module
- This is a natural join operation
 - But because we're practicing, we're going to use a standard CROSS JOIN and WHERE clause
 - Exercise: repeat the query using natural join

The Query so Far

```
SELECT some-information  
  FROM Student, Module, Grade  
  WHERE Student.ID = Grade.ID  
  AND Module.Code = Grade.Code  
  AND student-is-in-third-year
```

UNION

```
SELECT some-information  
  FROM Student, Module, Grade  
  WHERE Student.ID = Grade.ID  
  AND Module.Code = Grade.Code  
  AND student-is-in-first-or-second-year;
```

Information for Finalists

- We must retrieve
 - Computed average mark, weighted 40-60 across years 2 and 3
 - First year marks must be ignored
 - The ID, Name and Year are needed as they are used for ordering
- The average is difficult
 - We don't have any statements to separate years 2 and 3 easily
 - We can exploit the fact that $40 = 20 * 2$ and $60 = 20 * 3$, so YearTaken and the weighting have the same relationship

The Query so Far

```
SELECT some-information
  FROM Student, Module, Grade
 WHERE Student.ID = Grade.ID
  AND Module.Code = Grade.Code
  AND student-is-in-third-year
UNION
...
```

Information for Finalists

```
SELECT Year, Student.ID, Last, First,  
        SUM( ( (20*YearTaken) / 100) * Mark * Credits) /  
120  
        AS AverageMark  
FROM  
        Student, Module, Grade  
WHERE  
        Student.ID = Grade.ID  
        AND  
        Module.Code = Grade.Code  
        AND  
        YearTaken IN (2,3)  
        AND  
        Year = 3  
GROUP BY  
        Year, Student.ID, First, Last
```

Information for Others

- Other students are easier than finalists
 - We just need their average marks where YearTaken and Year are the same
 - As before, we need ID, Name and Year for ordering

. . .

UNION

```
SELECT some-information
FROM Student, Module, Grade
WHERE Student.ID = Grade.ID
AND Module.Code = Grade.Code
AND student-is-in-first-or-second-
year;
```

Information for Finalists

```
SELECT Year, Student.ID, Last, First,  
        SUM(Mark*Credits)/120 AS AverageMark  
FROM  
        Student, Module, Grade  
WHERE  
        Student.ID = Grade.ID  
AND  
        Module.Code = Grade.Code  
AND  
        YearTaken = Year  
AND  
        Year IN (1,2)  
GROUP BY  
        Year, Student.ID, First, Last
```

```

SELECT Year, Student.ID, Last, First,
        SUM(((20*YearTaken)/100)*Mark*Credits)/120
        AS AverageMark
FROM Student, Module, Grade
WHERE Student.ID = Grade.ID
        AND Module.Code = Grade.Code
        AND YearTaken IN (2,3)
        AND Year = 3
GROUP BY Year, Student.ID, Last, First

UNION

SELECT Year, Student.ID, Last, First,
        SUM(Mark*Credits)/120 AS AverageMark
FROM Student, Module, Grade
WHERE Student.ID = Grade.ID
        AND Module.Code = Grade.Code
        AND YearTaken = Year
        AND Year IN (1,2)
GROUP BY Year, Student.ID, Last, First

ORDER BY
        Year desc, AverageMark desc, Last, First, ID;

```



Questions?